

Regression

Lev Mazaev

Exercise 1 (Applied Predictive Modeling, p. 137)

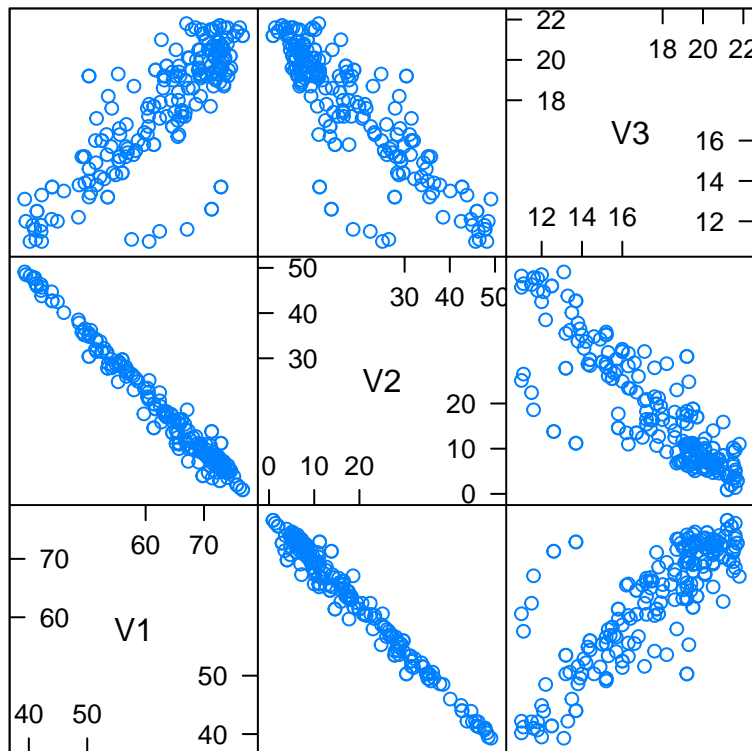
Answers to the questions are in the end of the document.

Loading packages and data

```
library(caret)
library(corrplot)
library(doMC)
registerDoMC(8)
data(tecator)
```

Checking the data

```
splom(~endpoints)
```



Scatter Plot Matrix

```
any(is.na(absorp))
```

```
## [1] FALSE
```

```
any(is.na(endpoints))

## [1] FALSE

colnames(absorp) <- paste0("V", 1:100)
head(absorp[, 1:6])

##           V1          V2          V3          V4          V5          V6
## [1,] 2.61776 2.61814 2.61859 2.61912 2.61981 2.62071
## [2,] 2.83454 2.83871 2.84283 2.84705 2.85138 2.85587
## [3,] 2.58284 2.58458 2.58629 2.58808 2.58996 2.59192
## [4,] 2.82286 2.82460 2.82630 2.82814 2.83001 2.83192
## [5,] 2.78813 2.78989 2.79167 2.79350 2.79538 2.79746
## [6,] 3.00993 3.01540 3.02086 3.02634 3.03190 3.03756

head(endpoints)

##           [,1] [,2] [,3]
## [1,] 60.5 22.5 16.7
## [2,] 46.0 40.1 13.5
## [3,] 71.0 8.4 20.5
## [4,] 72.8 5.9 20.7
## [5,] 58.3 25.5 15.5
## [6,] 44.0 42.7 13.7
```

Plotting 10 random spectra

```
p10randspectra <- function() {
  set.seed(1)
  inSubset <- sample(1:dim(endpoints)[1], 10)

  absorpSubset <- absorp[inSubset,]
  endpointSubset <- endpoints[inSubset, 3]

  newOrder <- order(absorpSubset[,1])
  absorpSubset <- absorpSubset[newOrder,]
  endpointSubset <- endpointSubset[newOrder]


  plotColors <- rainbow(10)

  plot(absorpSubset[1,],
       type = "n",
       ylim = range(absorpSubset),
       xlim = c(0, 105),
       xlab = "Wavelength Index",
       ylab = "Absorption")

  for(i in 1:10)
  {
    points(absorpSubset[i,], type = "l", col = plotColors[i], lwd = 2)
    text(105, absorpSubset[i,100], endpointSubset[i], col = plotColors[i])
  }
  title("Predictor Profiles for 10 Random Samples")
}
```

```
}  
p10randspectra()
```


Predictor Profiles for 10 Random Samples



Correlation matrix

All predictors are very highly correlated

```
corrplot(corr = cor(absorp), order = "hclust")
```



Principal Component Analysis


```
pcaObject <- prcomp(absorp, center = TRUE, scale. = TRUE)
summary(pcaObject)
```

```
## Importance of components:
##              PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation  9.9311  0.9847  0.52851  0.33827  0.08038  0.05123
## Proportion of Variance 0.9863  0.0097  0.00279  0.00114  0.00006  0.00003
## Cumulative Proportion 0.9863  0.9960  0.99875  0.99990  0.99996  0.99999
##              PC7      PC8      PC9      PC10     PC11     PC12
## Standard deviation  0.02681  0.01961  0.008564  0.006739  0.004442  0.003361
## Proportion of Variance 0.00001  0.00000  0.000000  0.000000  0.000000  0.000000
## Cumulative Proportion 0.99999  1.00000  1.000000  1.000000  1.000000  1.000000
##              PC13     PC14     PC15     PC16     PC17
## Standard deviation  0.001867  0.001377  0.0009449  0.0008641  0.0007558
## Proportion of Variance 0.000000  0.000000  0.0000000  0.0000000  0.0000000
## Cumulative Proportion 1.000000  1.000000  1.0000000  1.0000000  1.0000000
##              PC18     PC19     PC20     PC21     PC22
## Standard deviation  0.0006977  0.0005884  0.0004628  0.0003897  0.0003341
## Proportion of Variance 0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
## Cumulative Proportion 1.0000000  1.0000000  1.0000000  1.0000000  1.0000000
##              PC23     PC24     PC25     PC26     PC27
## Standard deviation  0.0003123  0.0002721  0.0002616  0.000211  0.0001954
## Proportion of Variance 0.0000000  0.0000000  0.0000000  0.000000  0.0000000
```

## Cumulative Proportion	1.0000000	1.0000000	1.0000000	1.0000000	1.0000000
##	PC28	PC29	PC30	PC31	PC32
## Standard deviation	0.0001857	0.0001729	0.0001656	0.0001539	0.0001473
## Proportion of Variance	0.0000000	0.0000000	0.0000000	0.0000000	0.0000000
## Cumulative Proportion	1.0000000	1.0000000	1.0000000	1.0000000	1.0000000
##	PC33	PC34	PC35	PC36	PC37
## Standard deviation	0.0001392	0.0001339	0.0001269	0.0001082	0.000104
## Proportion of Variance	0.0000000	0.0000000	0.0000000	0.0000000	0.0000000
## Cumulative Proportion	1.0000000	1.0000000	1.0000000	1.0000000	1.0000000
##	PC38	PC39	PC40	PC41	PC42
## Standard deviation	9.98e-05	9.081e-05	8.668e-05	8.026e-05	7.762e-05
## Proportion of Variance	0.00e+00	0.000e+00	0.000e+00	0.000e+00	0.000e+00
## Cumulative Proportion	1.00e+00	1.000e+00	1.000e+00	1.000e+00	1.000e+00
##	PC43	PC44	PC45	PC46	PC47
## Standard deviation	7.36e-05	6.808e-05	6.541e-05	6.44e-05	5.897e-05
## Proportion of Variance	0.00e+00	0.000e+00	0.000e+00	0.00e+00	0.000e+00
## Cumulative Proportion	1.00e+00	1.000e+00	1.000e+00	1.00e+00	1.000e+00
##	PC48	PC49	PC50	PC51	PC52
## Standard deviation	5.422e-05	5.027e-05	4.893e-05	4.608e-05	4.419e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.000e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.000e+00	1.000e+00
##	PC53	PC54	PC55	PC56	PC57
## Standard deviation	4.037e-05	3.854e-05	3.8e-05	3.64e-05	3.497e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.0e+00	0.00e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.0e+00	1.00e+00	1.000e+00
##	PC58	PC59	PC60	PC61	PC62
## Standard deviation	3.443e-05	3.264e-05	3.104e-05	3.04e-05	2.959e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.00e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.00e+00	1.000e+00
##	PC63	PC64	PC65	PC66	PC67
## Standard deviation	2.844e-05	2.699e-05	2.586e-05	2.388e-05	2.364e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.000e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.000e+00	1.000e+00
##	PC68	PC69	PC70	PC71	PC72
## Standard deviation	2.284e-05	2.173e-05	2.058e-05	1.997e-05	1.93e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.000e+00	0.00e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.000e+00	1.00e+00
##	PC73	PC74	PC75	PC76	PC77
## Standard deviation	1.854e-05	1.807e-05	1.728e-05	1.693e-05	1.612e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.000e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.000e+00	1.000e+00
##	PC78	PC79	PC80	PC81	PC82
## Standard deviation	1.569e-05	1.516e-05	1.445e-05	1.408e-05	1.356e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.000e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.000e+00	1.000e+00
##	PC83	PC84	PC85	PC86	PC87
## Standard deviation	1.275e-05	1.224e-05	1.178e-05	1.09e-05	1.045e-05
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.00e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.00e+00	1.000e+00
##	PC88	PC89	PC90	PC91	PC92
## Standard deviation	1.009e-05	9.396e-06	8.728e-06	8.27e-06	7.613e-06
## Proportion of Variance	0.000e+00	0.000e+00	0.000e+00	0.00e+00	0.000e+00
## Cumulative Proportion	1.000e+00	1.000e+00	1.000e+00	1.00e+00	1.000e+00
##	PC93	PC94	PC95	PC96	PC97

```
## Standard deviation      6.83e-06 6.383e-06 5.946e-06 5.478e-06 4.826e-06
## Proportion of Variance 0.00e+00 0.000e+00 0.000e+00 0.000e+00 0.000e+00
## Cumulative Proportion  1.00e+00 1.000e+00 1.000e+00 1.000e+00 1.000e+00
##                        PC98      PC99      PC100
## Standard deviation      4.521e-06 4.164e-06 4.122e-06
## Proportion of Variance  0.000e+00 0.000e+00 0.000e+00
## Cumulative Proportion  1.000e+00 1.000e+00 1.000e+00
```

```
plot(1:100, 100*summary(pcaObject)$importance[2, ], type = "b",
     pch = 19, xlab = "# of component", ylab = "percent of variance")
```



```
rm(pcaObject)
```

The effective dimension of the data is 1 because PC1 catches 98.6% of variance

```
PCAFit <- preprocess(absorp, method = c("pca"))
PCAFit
```

```
## Created from 215 samples and 100 variables
##
## Pre-processing:
## - centered (100)
## - ignored (0)
## - principal component signal extraction (100)
## - scaled (100)
##
## PCA needed 2 components to capture 95 percent of the variance
```

```
rm(PCAFit)
```


PCA from preprocess chooses 2 components

Pre-processing and splitting the data

```
trainrows <- createDataPartition(1:215, p = 0.6, list = TRUE)[[1]]
transFit <- preprocess(absorp, method = c("BoxCox", "center", "scale"))
transAbsorp <- as.data.frame(predict(transFit, absorp))
trainX <- transAbsorp[trainrows, ]
testX <- transAbsorp[-trainrows, ]
transEndpoints <- as.data.frame(endpoints)
trainY <- transEndpoints[trainrows, 2] # col #2 because only fat will be predicted
testY <- transEndpoints[-trainrows, 2]
trainingData <- cbind(trainX, trainY)
colnames(trainingData)[101] <- "FatP"
```

Visualization

```
ggplot(trainingData, aes(x = V1, y = FatP)) +  
  geom_point()
```



Function to evaluate models (by defaultSummary)

```
modeval <- function(model, X = testX, Y = testY) {  
  pred <- predict(model, X)  
  values <- data.frame(obs = Y, pred = pred)  
}
```

```

return(defaultSummary(values))
}

```

Ordinary Linear Regression, lm function

```

lmbasic <- lm(FatP ~ ., data = trainingData)
summary(lmbasic)

```

```

##
## Call:
## lm(formula = FatP ~ ., data = trainingData)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.91560 -0.24840 -0.05153  0.26542  1.29998
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  1.828e+01  1.497e-01 122.056 < 2e-16 ***
## V1          -4.629e+02  1.548e+03  -0.299  0.766928
## V2          -1.942e+03  3.073e+03  -0.632  0.532207
## V3           1.014e+02  5.743e+03   0.018  0.986027
## V4           6.734e+03  9.207e+03   0.731  0.470211
## V5          -7.771e+03  1.084e+04  -0.717  0.478914
## V6           1.000e+04  9.226e+03   1.084  0.286833
## V7          -8.322e+03  5.288e+03  -1.574  0.126002
## V8           2.490e+03  3.673e+03   0.678  0.502943
## V9          -2.486e+02  3.045e+03  -0.082  0.935476
## V10          -2.583e+03  3.483e+03  -0.742  0.464092
## V11           8.111e+03  4.803e+03   1.689  0.101614
## V12          -1.025e+04  6.400e+03  -1.602  0.119595
## V13           1.307e+02  4.651e+03   0.028  0.977777
## V14          -8.119e+02  5.851e+03  -0.139  0.890573
## V15           7.559e+03  8.899e+03   0.849  0.402392
## V16          -4.688e+03  6.067e+03  -0.773  0.445755
## V17           4.404e+03  3.445e+03   1.278  0.210936
## V18           3.271e+02  3.928e+03   0.083  0.934185
## V19           7.619e+01  4.748e+03   0.016  0.987304
## V20          -3.289e+02  5.512e+03  -0.060  0.952818
## V21          -1.817e+04  6.598e+03  -2.753  0.009916 **
## V22           3.715e+04  9.633e+03   3.857  0.000565 ***
## V23          -3.531e+04  9.469e+03  -3.729  0.000799 ***
## V24           1.313e+04  6.903e+03   1.903  0.066708 .
## V25           1.519e+03  5.302e+03   0.287  0.776406
## V26           3.830e+03  4.064e+03   0.942  0.353512
## V27          -4.221e+03  4.224e+03  -0.999  0.325666
## V28          -3.479e+03  4.836e+03  -0.719  0.477428
## V29           7.589e+03  5.117e+03   1.483  0.148438
## V30          -6.984e+03  5.955e+03  -1.173  0.250090
## V31           4.768e+03  1.063e+04   0.448  0.657036
## V32          -1.622e+03  1.253e+04  -0.129  0.897879
## V33          -1.596e+01  9.312e+03  -0.002  0.998644

```



## V34	-4.273e+03	6.644e+03	-0.643	0.524968
## V35	3.266e+03	4.934e+03	0.662	0.513081
## V36	-3.711e+03	3.978e+03	-0.933	0.358339
## V37	2.642e+03	4.952e+03	0.533	0.597662
## V38	5.443e+03	6.558e+03	0.830	0.413123
## V39	-5.527e+03	8.271e+03	-0.668	0.509107
## V40	5.475e+01	1.056e+04	0.005	0.995898
## V41	2.243e+03	1.105e+04	0.203	0.840475
## V42	3.979e+03	9.709e+03	0.410	0.684874
## V43	-4.950e+03	5.932e+03	-0.834	0.410653
## V44	-5.449e+03	2.581e+03	-2.111	0.043211 *
## V45	1.088e+04	5.688e+03	1.913	0.065349 .
## V46	-7.409e+03	3.891e+03	-1.904	0.066501 .
## V47	3.031e+03	3.494e+03	0.868	0.392518
## V48	-1.012e+03	3.440e+03	-0.294	0.770722
## V49	-2.981e+03	4.001e+03	-0.745	0.462069
## V50	4.557e+03	5.327e+03	0.855	0.399107
## V51	-1.681e+02	5.243e+03	-0.032	0.974628
## V52	6.316e+02	4.197e+03	0.151	0.881376
## V53	-6.221e+03	3.551e+03	-1.752	0.090046 .
## V54	2.175e+03	3.860e+03	0.564	0.577210
## V55	9.994e+03	6.833e+03	1.463	0.153984
## V56	-1.185e+04	6.635e+03	-1.786	0.084257 .
## V57	2.176e+03	4.272e+03	0.509	0.614152
## V58	2.910e+03	3.049e+03	0.955	0.347458
## V59	-1.883e+03	2.393e+03	-0.787	0.437647
## V60	1.373e+03	2.123e+03	0.647	0.522782
## V61	-2.817e+03	2.888e+03	-0.975	0.337269
## V62	7.674e+03	3.709e+03	2.069	0.047282 *
## V63	-4.868e+03	4.459e+03	-1.092	0.283642
## V64	-4.672e+03	6.178e+03	-0.756	0.455434
## V65	7.054e+03	7.344e+03	0.960	0.344500
## V66	-2.219e+03	7.632e+03	-0.291	0.773231
## V67	-8.189e+03	8.457e+03	-0.968	0.340651
## V68	1.320e+04	8.196e+03	1.611	0.117665
## V69	-5.371e+03	6.807e+03	-0.789	0.436280
## V70	-3.230e+03	6.138e+03	-0.526	0.602614
## V71	5.738e+02	5.247e+03	0.109	0.913660
## V72	4.002e+03	4.147e+03	0.965	0.342322
## V73	3.961e+02	4.059e+03	0.098	0.922907
## V74	-5.646e+03	3.939e+03	-1.433	0.162080
## V75	6.331e+03	3.569e+03	1.774	0.086256 .
## V76	-1.783e+03	2.659e+03	-0.671	0.507642
## V77	-1.472e+03	3.114e+03	-0.473	0.639844
## V78	5.584e+03	4.082e+03	1.368	0.181494
## V79	-2.909e+03	4.033e+03	-0.721	0.476233
## V80	-4.971e+03	4.552e+03	-1.092	0.283530
## V81	7.322e+02	6.194e+03	0.118	0.906677
## V82	-2.500e+03	6.461e+03	-0.387	0.701508
## V83	9.551e+03	6.330e+03	1.509	0.141824
## V84	-2.057e+04	1.086e+04	-1.894	0.067901 .
## V85	2.190e+04	1.203e+04	1.820	0.078754 .
## V86	-2.656e+03	9.613e+03	-0.276	0.784211
## V87	-8.226e+03	8.442e+03	-0.974	0.337682

```
## V88      6.071e+03  7.268e+03   0.835 0.410133
## V89      3.563e+03  7.654e+03   0.466 0.644921
## V90     -9.866e+03  7.737e+03  -1.275 0.212038
## V91      7.244e+03  8.206e+03   0.883 0.384398
## V92      1.234e+03  9.512e+03   0.130 0.897651
## V93     -7.118e+03  9.387e+03  -0.758 0.454197
## V94      8.136e+03  9.444e+03   0.861 0.395809
## V95     -4.353e+03  5.011e+03  -0.869 0.391898
## V96     -1.561e+03  3.782e+03  -0.413 0.682691
## V97      1.169e+03  5.407e+03   0.216 0.830269
## V98     -2.322e+03  4.942e+03  -0.470 0.641867
## V99      9.215e+02  5.592e+03   0.165 0.870223
## V100     1.345e+03  2.353e+03   0.572 0.571628
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.847 on 30 degrees of freedom
## Multiple R-squared:  0.9989, Adjusted R-squared:  0.9952
## F-statistic: 270.8 on 100 and 30 DF,  p-value: < 2.2e-16
```

```
modeval(lmbasic)
```

```
##      RMSE  Rsquared      MAE
## 2.9617071 0.9519516 2.0495377
```

```
modstats <- t(as.data.frame(modeval(lmbasic)))
rownames(modstats)[1] <- "LM basic R"
plot(testY, predict(lmbasic, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```



Ordinary Linear Regression, lm from caret


```
ctrl <- trainControl(method = "cv", number = 10)
lmFit <- train(x = trainX, y = trainY, method = "lm", trControl = ctrl)
lmFit
```

```
## Linear Regression
##
## 131 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 118, 118, 116, 118, 119, 118, ...
## Resampling results:
##
##      RMSE      Rsquared   MAE
##  4.893908  0.8749618  3.501878
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
```

```
modeval(lmFit)
```

```
##      RMSE Rsquared   MAE
## 2.9617071 0.9519516 2.0495377
```

```
modstats <- rbind(modstats, modeval(lmFit))
rownames(modstats)[2] <- "LM caret"
plot(testY, predict(lmFit, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```




Ordinary Linear Regression with filtering of highly correlated values

```
tooHigh <- findCorrelation(cor(trainX), cutoff = 0.99)
length(tooHigh)
```

```
## [1] 98
trainXfiltered <- trainX[, -tooHigh]
testXfiltered <- testX[, -tooHigh]
lmFiltered <- train(x = trainXfiltered, y = trainY,
                   method = "lm", trControl = ctrl)
lmFiltered

## Linear Regression
##
## 131 samples
## 2 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 119, 119, 117, 118, 117, 118, ...
## Resampling results:
##
## RMSE      Rsquared    MAE
## 9.557742  0.4272833  7.737826
##
## Tuning parameter 'intercept' was held constant at a value of TRUE
modeval(lmFiltered, X = testXfiltered)

## RMSE      Rsquared    MAE
## 10.3175997 0.4119963  8.0231636
modstats <- rbind(modstats, modeval(lmFiltered, X = testXfiltered))
rownames(modstats)[3] <- "LM filtered"
plot(testY, predict(lmFiltered, testXfiltered), pch = 19, xlab = "Observed", ylab = "Predicted")
```



Robust Linear Regression, rlm from MASS

```
library(MASS)
rlmFit <- rlm(FatP ~ ., data = trainingData)

## Warning in rlm.default(x, y, weights, method = method, wt.method =
## wt.method, : 'rlm' failed to converge in 20 steps

summary(rlmFit)

##
## Call: rlm(formula = FatP ~ ., data = trainingData)
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2.4619189 -0.0199005 -0.0003726  0.0181482  3.3265732
##
## Coefficients:
##      Value      Std. Error t value
## (Intercept)    18.3761      0.0138 1335.0451
## V1             2770.6347    142.2602   19.4758
## V2            -7298.8485    282.5009  -25.8366
## V3             3063.5439    527.9737    5.8025
## V4             5566.3191    846.3872    6.5766
## V5           -10883.4683    996.3101  -10.9238
## V6            13669.5039    848.0823   16.1181
## V7            -9272.4878    486.0814  -19.0760
## V8             4231.2882    337.6115   12.5330
## V9             435.6881    279.9365    1.5564
## V10           -3262.2660    320.1872  -10.1886
## V11            4990.9074    441.4776   11.3050
## V12           -6964.9200    588.3377  -11.8383
## V13           -2548.5699    427.5765   -5.9605
## V14           -1570.7126    537.8888   -2.9201
## V15            13268.9496    818.0370   16.2205
## V16           -9029.0281    557.7159  -16.1893
## V17            5507.6272    316.7185   17.3897
## V18            2552.8091    361.0687    7.0701
## V19           -5270.2011    436.5035  -12.0737
## V20            3281.4370    506.7061    6.4760
## V21           -15316.1554    606.5358  -25.2519
## V22            26730.6917    885.4885   30.1875
## V23           -24734.8057    870.4901  -28.4148
## V24            8049.0797    634.5428   12.6848
## V25            4914.6899    487.3985   10.0835
## V26            2153.3341    373.5707    5.7642
## V27           -5686.2240    388.3103  -14.6435
## V28           -3822.6584    444.5125   -8.5997
## V29            10318.0984    470.3602   21.9366
## V30           -6896.4375    547.3866  -12.5988
## V31            4311.6326    977.2368    4.4121
## V32           -6855.7322   1152.0293   -5.9510
## V33            6733.4215    856.0403    7.8658
## V34           -6030.1009    610.7172   -9.8738
## V35             393.3382    453.5466    0.8672
```


## V36	1064.8944	365.6967	2.9120
## V37	244.1951	455.2480	0.5364
## V38	3614.6406	602.8809	5.9956
## V39	-6752.2724	760.2826	-8.8813
## V40	7375.1906	970.8101	7.5969
## V41	-3230.7212	1015.6892	-3.1808
## V42	4344.0515	892.5309	4.8671
## V43	-3894.0512	545.3409	-7.1406
## V44	-6076.1117	237.2634	-25.6091
## V45	11252.3001	522.8719	21.5202
## V46	-8577.8163	357.6493	-23.9839
## V47	5207.7855	321.1987	16.2136
## V48	-1470.1216	316.2524	-4.6486
## V49	-3823.0969	367.7824	-10.3950
## V50	2521.4502	489.6783	5.1492
## V51	3537.4867	481.9563	7.3399
## V52	-987.2358	385.8046	-2.5589
## V53	-6137.7572	326.4689	-18.8004
## V54	3838.7904	354.8047	10.8194
## V55	5912.6381	628.1547	9.4127
## V56	-9554.4487	609.9312	-15.6648
## V57	4701.4954	392.6766	11.9729
## V58	-583.2631	280.3006	-2.0808
## V59	-926.8589	220.0126	-4.2128
## V60	2584.4644	195.1934	13.2405
## V61	-4769.0360	265.5223	-17.9610
## V62	4901.9000	340.9797	14.3759
## V63	3677.4417	409.9242	8.9710
## V64	-9375.2076	567.9233	-16.5079
## V65	4869.9198	675.1466	7.2131
## V66	-1837.3990	701.6072	-2.6188
## V67	-2566.4278	777.4626	-3.3010
## V68	8768.4660	753.4153	11.6383
## V69	-6659.1568	625.7213	-10.6424
## V70	-81.0452	564.2104	-0.1436
## V71	-377.6398	482.3735	-0.7829
## V72	4121.4926	381.2608	10.8102
## V73	-5490.7170	373.1020	-14.7164
## V74	1087.3778	362.0761	3.0032
## V75	5755.8616	328.1296	17.5414
## V76	-2967.4297	244.4496	-12.1392
## V77	168.9009	286.2800	0.5900
## V78	6493.9110	375.2388	17.3061
## V79	-6141.4021	370.6981	-16.5671
## V80	-5107.9144	418.4830	-12.2058
## V81	1953.1412	569.3530	3.4305
## V82	-2180.4742	593.9336	-3.6712
## V83	5140.4764	581.9129	8.8338
## V84	-12238.7120	998.1301	-12.2616
## V85	17668.0528	1105.9253	15.9758
## V86	-2842.9511	883.7185	-3.2170
## V87	-6406.6042	776.0650	-8.2552
## V88	2324.6917	668.0762	3.4797
## V89	6392.2113	703.6409	9.0845

```
## V90      -9756.6156    711.2069   -13.7184
## V91      -1402.1908    754.3505    -1.8588
## V92      17930.1291    874.3721    20.5063
## V93     -20289.8542    862.9119   -23.5132
## V94      13695.8545    868.1389    15.7761
## V95     -4763.8510    460.6303   -10.3420
## V96     -4584.4534    347.6886   -13.1855
## V97       6688.3539    497.0259    13.4568
## V98     -10357.9420    454.2568   -22.8020
## V99       7525.8916    514.0890    14.6393
## V100     -643.7613    216.2599    -2.9768
##
## Residual standard error: 0.03033 on 30 degrees of freedom
```

```
modeval(rlmFit)
```

```
##      RMSE  Rsquared      MAE
## 3.6285964 0.9316943 2.4430944
```

```
modstats <- rbind(modstats, modeval(rlmFit))
rownames(modstats)[4] <- "RLM MASS"
plot(testY, predict(rlmFit, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```



Robust Linear Regression with PCA

```
rlmPCA <- train(x = trainX, y = trainY, method = "rlm", preProcess = "pca",
               trControl = ctrl)
summary(rlmPCA)
```


```
##
## Call: rlm(formula = .outcome ~ ., data = dat, psi = psi)
## Residuals:
##      Min       1Q   Median       3Q      Max
```

```
## -17.800 -6.627 -1.751 6.751 30.745
##
## Coefficients:
##          Value Std. Error t value
## (Intercept) 16.1510 0.8912 18.1232
## PC1          0.5524 0.0901 6.1334
## PC2          3.0393 0.9011 3.3728
##
## Residual standard error: 9.871 on 128 degrees of freedom
```

```
modeval(rlmPCA)
```

```
##          RMSE  Rsquared      MAE
## 11.5720269 0.2833121 9.1018698
```

```
modstats <- rbind(modstats, modeval(rlmPCA))
rownames(modstats)[5] <- "RLM PCA"
plot(testY, predict(rlmPCA, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```



Partial Least Squares Regression

```
indx <- createFolds(trainY, returnTrain = TRUE)
ctrl <- trainControl(method = "cv", index = indx)
plsTune <- train(x = trainX, y = trainY, method = "pls",
  tuneGrid = expand.grid(ncomp = 1:50),
  trControl = ctrl)
```

Tuning parameters

```
plsTune
```

```
## Partial Least Squares
##
```




```

## 131 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 117, 117, 118, 118, 118, 118, ...
## Resampling results across tuning parameters:
##
##   ncomp  RMSE      Rsquared  MAE
##   1      10.797966  0.2310967  8.7416942
##   2       8.192915  0.5648930  6.7464233
##   3       5.240568  0.8104156  4.1004251
##   4       4.624434  0.8642693  3.6080863
##   5       3.059443  0.9488733  2.2870680
##   6       2.986761  0.9541787  2.2186577
##   7       2.815218  0.9611686  2.0701101
##   8       2.495234  0.9710269  1.7576628
##   9       2.374636  0.9706905  1.7333109
##  10       2.099915  0.9763563  1.5480682
##  11       2.046205  0.9800377  1.4812418
##  12       1.956509  0.9816527  1.4097997
##  13       1.734992  0.9839708  1.3243399
##  14       1.869780  0.9812099  1.3514968
##  15       1.672174  0.9844472  1.2201141
##  16       1.487821  0.9866369  1.0984136
##  17       1.428482  0.9878811  1.0601455
##  18       1.342269  0.9887071  1.0152220
##  19       1.344269  0.9887766  1.0142185
##  20       1.330680  0.9890851  1.0110143
##  21       1.277281  0.9891429  0.9573668
##  22       1.256334  0.9896627  0.9510145
##  23       1.236562  0.9899599  0.9421029
##  24       1.244488  0.9900481  0.9461620
##  25       1.267352  0.9895641  0.9552685
##  26       1.336248  0.9884864  1.0056685
##  27       1.396058  0.9875957  1.0332871
##  28       1.470713  0.9860245  1.0756039
##  29       1.554574  0.9843602  1.1072995
##  30       1.575198  0.9840256  1.1301841
##  31       1.622658  0.9827225  1.1530157
##  32       1.649369  0.9821588  1.1810998
##  33       1.718049  0.9806263  1.2260404
##  34       1.762263  0.9794269  1.2455132
##  35       1.844402  0.9775971  1.2945704
##  36       1.873360  0.9770406  1.3258212
##  37       1.943670  0.9753431  1.3694065
##  38       1.993889  0.9737307  1.3899052
##  39       2.057341  0.9723603  1.4283037
##  40       2.153605  0.9700275  1.5060714
##  41       2.289832  0.9661782  1.6045963
##  42       2.416229  0.9619112  1.6876382
##  43       2.537148  0.9582574  1.7569513
##  44       2.634238  0.9550149  1.8044361
##  45       2.706725  0.9521107  1.8319188

```

```
## 46      2.798918  0.9480944  1.8584917
## 47      2.878299  0.9451495  1.8938869
## 48      2.940079  0.9425120  1.9341781
## 49      2.992495  0.9402794  1.9636487
## 50      3.046340  0.9382202  2.0011228
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 23.
```


```
plot(plsTune)
```



```
modeval(plsTune)
```

```
##      RMSE Rsquared      MAE
## 1.450754 0.987997 1.059324
```

```
modstats <- rbind(modstats, modeval(plsTune))
rownames(modstats)[6] <- "PLSR"
plot(testY, predict(plsTune, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```



Principal Component Regression

```
pcrTune <- train(x = trainX, y = trainY, method = "pcr",
  tuneGrid = expand.grid(ncomp = 1:50),
  trControl = ctrl)
```

Tuning parameters


```
pcrTune
```

```
## Principal Component Analysis
##
## 131 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 117, 117, 118, 118, 118, 118, ...
## Resampling results across tuning parameters:
##
##  ncomp  RMSE      Rsquared  MAE
##    1    10.820029  0.2279488  8.7661806
##    2    10.395319  0.2848395  8.2763549
##    3     7.701887  0.6016426  6.1666383
##    4     4.935712  0.8413751  3.8252846
##    5     3.082908  0.9478825  2.3047073
##    6     3.119444  0.9461726  2.3485515
##    7     3.206411  0.9442564  2.3917926
##    8     2.937604  0.9559360  2.1733786
##    9     3.014202  0.9550826  2.1952601
##   10     2.796865  0.9616962  2.0389795
```

```

## 11      2.676876  0.9644601  1.9022231
## 12      2.771657  0.9602406  1.9632980
## 13      2.062902  0.9799680  1.5039108
## 14      2.007329  0.9798668  1.4465465
## 15      2.032164  0.9798037  1.4509418
## 16      1.979465  0.9810269  1.4438343
## 17      2.005106  0.9807589  1.4493639
## 18      1.850263  0.9832113  1.3330298
## 19      1.771701  0.9822700  1.2734702
## 20      1.842635  0.9813318  1.3081510
## 21      1.671934  0.9836261  1.2024552
## 22      1.553051  0.9860597  1.1803341
## 23      1.340847  0.9887691  1.0171450
## 24      1.366147  0.9883475  1.0344072
## 25      1.378025  0.9881334  1.0403588
## 26      1.377411  0.9882482  1.0333659
## 27      1.379409  0.9882609  1.0366295
## 28      1.386430  0.9880385  1.0379671
## 29      1.409747  0.9874717  1.0635313
## 30      1.438352  0.9872350  1.0868188
## 31      1.387631  0.9878453  1.0789480
## 32      1.345721  0.9885926  1.0680052
## 33      1.295375  0.9891948  1.0293526
## 34      1.310491  0.9890261  1.0280200
## 35      1.277249  0.9896162  0.9789182
## 36      1.283031  0.9896241  0.9702687
## 37      1.288451  0.9893819  0.9548171
## 38      1.273877  0.9896686  0.9359205
## 39      1.296278  0.9893254  0.9612144
## 40      1.270711  0.9900889  0.9431444
## 41      1.327426  0.9892806  0.9761604
## 42      1.322145  0.9893100  0.9861686
## 43      1.315049  0.9893868  0.9790014
## 44      1.362870  0.9889922  1.0194829
## 45      1.408792  0.9883322  1.0435443
## 46      1.407356  0.9883655  1.0419257
## 47      1.434110  0.9876521  1.0659104
## 48      1.387823  0.9883715  1.0410124
## 49      1.348606  0.9885151  1.0109201
## 50      1.357893  0.9883561  1.0157296
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was ncomp = 40.
plot(pcrTune)

```




```
modeval(pcrTune)
```

```
##      RMSE Rsquared      MAE
## 1.4114441 0.9885807 1.0174803
```

```
modstats <- rbind(modstats, modeval(pcrTune))
```

```
rownames(modstats)[7] <- "PCR"
```

```
plot(testY, predict(pcrTune, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```




Ridge Regression

```
ridgeGrid <- data.frame(.lambda = seq(0, .1, length = 15))
ridgeRegFit <- train(x = trainX, y = trainY, method = "ridge", tuneGrid = ridgeGrid,
                    trControl = ctrl)
```

Tuning parameters

ridgeRegFit

```
## Ridge Regression
##
## 131 samples
## 100 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 117, 117, 118, 118, 118, 118, ...
## Resampling results across tuning parameters:
##
##   lambda      RMSE      Rsquared    MAE
## 0.000000000  4.648676  0.8739361  3.053590
## 0.007142857  3.564346  0.9264257  2.704078
## 0.014285714  3.951852  0.9052381  2.977535
## 0.021428571  4.198426  0.8910699  3.121035
## 0.028571429  4.382477  0.8805028  3.234836
## 0.035714286  4.535012  0.8718656  3.342973
## 0.042857143  4.669545  0.8643297  3.443875
## 0.050000000  4.792486  0.8574560  3.531894
## 0.057142857  4.907102  0.8509979  3.616112
## 0.064285714  5.015175  0.8448108  3.707965
## 0.071428571  5.117760  0.8388070  3.792516
## 0.078571429  5.215540  0.8329326  3.871539
## 0.085714286  5.308993  0.8271537  3.947053
## 0.092857143  5.398487  0.8214493  4.020327
## 0.100000000  5.484324  0.8158064  4.092073
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was lambda = 0.007142857.
plot(ridgeRegFit)
```




```
modeval(ridgeRegFit)
```

```
##      RMSE  Rsquared    MAE
## 3.5626739 0.9278507 2.6659641
```

```
modstats <- rbind(modstats, modeval(ridgeRegFit))
```

```
rownames(modstats)[8] <- "Ridge"
```

```
plot(testY, predict(ridgeRegFit, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```



LASSO Regression

```
enetGrid <- expand.grid(.lambda = c(0, 0.01, .1),  
                        .fraction = seq(0.5, 1, length = 20))  
enetTune <- train(x = trainX, y = trainY, method = "enet", tuneGrid = enetGrid,  
                  trControl = ctrl)
```


Tuning parameters

enetTune

```
## Elasticnet  
##  
## 131 samples  
## 100 predictors  
##  
## No pre-processing  
## Resampling: Cross-Validated (10 fold)  
## Summary of sample sizes: 117, 117, 118, 118, 118, ...  
## Resampling results across tuning parameters:  
##  
##   lambda  fraction  RMSE      Rsquared  MAE  
##   0.00    0.5000000  2.655872  0.9553970  1.776683  
##   0.00    0.5263158  2.764026  0.9517389  1.838204  
##   0.00    0.5526316  2.874207  0.9479933  1.904889  
##   0.00    0.5789474  2.991536  0.9439906  1.978107  
##   0.00    0.6052632  3.097154  0.9404009  2.044172  
##   0.00    0.6315789  3.203109  0.9366824  2.107306  
##   0.00    0.6578947  3.313290  0.9327638  2.174620  
##   0.00    0.6842105  3.417302  0.9286165  2.240152  
##   0.00    0.7105263  3.526303  0.9241090  2.305848  
##   0.00    0.7368421  3.636393  0.9197645  2.373436  
##   0.00    0.7631579  3.751580  0.9151503  2.437882  
##   0.00    0.7894737  3.859141  0.9104812  2.496985  
##   0.00    0.8157895  3.962087  0.9059274  2.561530  
##   0.00    0.8421053  4.055855  0.9016803  2.630589  
##   0.00    0.8684211  4.155228  0.8971424  2.701036  
##   0.00    0.8947368  4.253113  0.8925933  2.770857  
##   0.00    0.9210526  4.352258  0.8879968  2.841411  
##   0.00    0.9473684  4.451800  0.8833440  2.912634  
##   0.00    0.9736842  4.552012  0.8785606  2.984549  
##   0.00    1.0000000  4.648676  0.8739361  3.053590  
##   0.01    0.5000000  5.057557  0.8596438  3.835138  
##   0.01    0.5263158  4.889250  0.8672786  3.680157  
##   0.01    0.5526316  4.735316  0.8739756  3.551369  
##   0.01    0.5789474  4.596967  0.8797555  3.437292  
##   0.01    0.6052632  4.474406  0.8848446  3.341581  
##   0.01    0.6315789  4.366714  0.8892359  3.255599  
##   0.01    0.6578947  4.273453  0.8930382  3.181319  
##   0.01    0.6842105  4.193409  0.8962758  3.123712  
##   0.01    0.7105263  4.120861  0.8993046  3.072212  
##   0.01    0.7368421  4.055950  0.9020514  3.027957  
##   0.01    0.7631579  4.004021  0.9042753  3.000367  
##   0.01    0.7894737  3.961440  0.9061097  2.977436
```




```
## 0.01 0.8157895 3.924519 0.9077530 2.955062
## 0.01 0.8421053 3.892900 0.9092024 2.934966
## 0.01 0.8684211 3.863021 0.9106925 2.915994
## 0.01 0.8947368 3.834561 0.9121567 2.897791
## 0.01 0.9210526 3.808118 0.9135219 2.880316
## 0.01 0.9473684 3.783534 0.9147905 2.863686
## 0.01 0.9736842 3.761562 0.9159236 2.849526
## 0.01 1.0000000 3.745081 0.9167936 2.839046
## 0.10 0.5000000 7.133142 0.7042959 5.607195
## 0.10 0.5263158 7.004225 0.7158628 5.486939
## 0.10 0.5526316 6.883820 0.7262456 5.374492
## 0.10 0.5789474 6.764364 0.7361912 5.264182
## 0.10 0.6052632 6.648373 0.7454467 5.158882
## 0.10 0.6315789 6.536535 0.7539657 5.055459
## 0.10 0.6578947 6.428407 0.7618546 4.953551
## 0.10 0.6842105 6.323749 0.7691552 4.854308
## 0.10 0.7105263 6.224720 0.7757464 4.763113
## 0.10 0.7368421 6.133017 0.7815685 4.678516
## 0.10 0.7631579 6.047820 0.7867325 4.598173
## 0.10 0.7894737 5.967188 0.7914492 4.520387
## 0.10 0.8157895 5.890551 0.7957659 4.447580
## 0.10 0.8421053 5.819292 0.7996186 4.384847
## 0.10 0.8684211 5.752000 0.8031288 4.328816
## 0.10 0.8947368 5.689656 0.8062664 4.276132
## 0.10 0.9210526 5.632556 0.8090275 4.226434
## 0.10 0.9473684 5.580039 0.8114849 4.179379
## 0.10 0.9736842 5.530842 0.8137461 4.134220
## 0.10 1.0000000 5.484324 0.8158064 4.092073
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were fraction = 0.5 and lambda = 0.
plot(enetTune)
```



```
modeval(enetTune)
```

```
##      RMSE Rsquared      MAE
## 2.1013149 0.9748699 1.5205947
```

```
modstats <- rbind(modstats, modeval(enetTune))
rownames(modstats)[9] <- "LASSO"
plot(testY, predict(enetTune, testX), pch = 19, xlab = "Observed", ylab = "Predicted")
```



Comparison of models

modstats

##		RMSE	Rsquared	MAE
##	LM basic R	2.961707	0.9519516	2.049538
##	LM caret	2.961707	0.9519516	2.049538
##	LM filtered	10.317600	0.4119963	8.023164
##	RLM MASS	3.628596	0.9316943	2.443094
##	RLM PCA	11.572027	0.2833121	9.101870
##	PLSR	1.450754	0.9879970	1.059324
##	PCR	1.411444	0.9885807	1.017480
##	Ridge	3.562674	0.9278507	2.665964
##	LASSO	2.101315	0.9748699	1.520595

Answers

- The effective dimension of the data according to the PCA test is 1, because PC1 catches 98.6% of variance.
- Evaluations and plots for tuning parameters of some models are aforementioned.
- Principal Component Regression has the best predictive ability: the lowest *RMSE* and the highest R^2 . Some models, such as RLM with PCA and LM with filtering of highly correlated predictors, are significantly worse than others. Possibly these models lose some valuable information contained in predictors that were removed due to high correlation or during PCA.
- I would use PCR or PLSR for predicting the fat content of a sample because they show the highest and very similar performance on the test data. LASSO may also be considered, all other models are outperformed by these ones.