

exercise72

Lev Mazaev


Exercise 7.2 (Applied Predictive Modeling, p. 169)

Loading the packages

```
library(mlbench)
library(caret)
library(doMC)
library(earth)
registerDoMC(8)
```

Creating the data


```
set.seed(200)
trainingData <- mlbench.friedman1(200, sd = 1)
trainingData$x <- data.frame(trainingData$x)
featurePlot(trainingData$x, trainingData$y)
```



```
testData <- mlbench.friedman1(5000, sd = 1)
testData$x <- data.frame(testData$x)
```

Correlation check

```
corrplot::corrplot(cor(trainingData$x))
```



```
findCorrelation(cor(trainingData$x), cutoff = 0.75)
```

```
## integer(0)
```

There are no highly-correlated predictors

Neural Network

```
nnetGrid <- expand.grid(.decay = c(0, 0.01, .1), .size = 1:10, .bag = FALSE)
nnetModel <- train(x = trainingData$x, y = trainingData$y,
  method = "avNNet", tuneGrid = nnetGrid,
  trControl = trainControl(method = "cv"),
  preProcess = c("center", "scale"),
  linout = TRUE, trace = TRUE,
  MaxNWts = 10 * (ncol(trainingData$x) + 1) + 10 + 1,
  maxit = 500)
nnetModel
```


```
## Model Averaged Neural Network
##
## 200 samples
```

```

## 10 predictors
##
## Pre-processing: centered (10), scaled (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...
## Resampling results across tuning parameters:
##
##   decay  size  RMSE      Rsquared  MAE
##   0.00    1   2.481454  0.7469747  1.975542
##   0.00    2   2.436424  0.7649192  1.930859
##   0.00    3   1.877814  0.8573051  1.476834
##   0.00    4   1.905918  0.8541738  1.529084
##   0.00    5   1.993979  0.8409575  1.576757
##   0.00    6   2.733189  0.7264896  2.029163
##   0.00    7   4.727163  0.5186760  2.950852
##   0.00    8   4.232154  0.6227455  2.886581
##   0.00    9   5.830266  0.4324491  3.391461
##   0.00   10   5.452456  0.5700175  2.969528
##   0.01    1   2.398640  0.7703358  1.871041
##   0.01    2   2.400588  0.7710578  1.888898
##   0.01    3   2.048665  0.8326338  1.617814
##   0.01    4   1.943547  0.8453374  1.564313
##   0.01    5   2.060612  0.8232218  1.568054
##   0.01    6   2.147995  0.8115418  1.721255
##   0.01    7   2.314310  0.7908798  1.836394
##   0.01    8   2.260879  0.7998637  1.815767
##   0.01    9   2.205192  0.8016452  1.794733
##   0.01   10   2.433887  0.7698170  1.895905
##   0.10    1   2.403712  0.7687049  1.869539
##   0.10    2   2.421932  0.7689830  1.888883
##   0.10    3   2.163172  0.8143900  1.700189
##   0.10    4   2.060149  0.8297096  1.670812
##   0.10    5   2.113895  0.8213375  1.700141
##   0.10    6   2.133011  0.8251511  1.710712
##   0.10    7   2.126111  0.8217688  1.685677
##   0.10    8   2.122994  0.8237300  1.706942
##   0.10    9   2.049278  0.8280368  1.611528
##   0.10   10   2.289593  0.7878086  1.811428
##
## Tuning parameter 'bag' was held constant at a value of FALSE
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were size = 3, decay = 0 and bag
## = FALSE.

```

```
plot(nnetModel)
```



Performance evaluation

```
nnetPred <- predict(nnetModel, newdata = testData$x)
modelstats <- postResample(pred = nnetPred, obs = testData$y)
postResample(pred = nnetPred, obs = testData$y)
```

```
##      RMSE  Rsquared      MAE
## 1.9263163 0.8528292 1.4630365
```


K Nearest Neighbors

```
knnModel <- train(x = trainingData$x, y = trainingData$y,
                  method = "knn", preProcess = c("center", "scale"),
                  tuneLength = 10)
knnModel
```

```
## k-Nearest Neighbors
##
## 200 samples
## 10 predictors
##
## Pre-processing: centered (10), scaled (10)
## Resampling: Bootstrapped (25 reps)
## Summary of sample sizes: 200, 200, 200, 200, 200, 200, ...
## Resampling results across tuning parameters:
##
## k  RMSE      Rsquared  MAE
##  5  1.9263163 0.8528292 1.4630365
```

```
##      5  3.612753  0.4847321  2.934855
##      7  3.441582  0.5351824  2.792389
##      9  3.385434  0.5592947  2.740026
##     11  3.332242  0.5864511  2.700934
##     13  3.315092  0.6029810  2.686746
##     15  3.298171  0.6224845  2.674603
##     17  3.293481  0.6341255  2.664912
##     19  3.296045  0.6483953  2.676133
##     21  3.310329  0.6539922  2.694153
##     23  3.334949  0.6560518  2.714964
##
## RMSE was used to select the optimal model using the smallest value.
## The final value used for the model was k = 17.
```

```
plot(knnModel)
```



Performance evaluation

```
knnPred <- predict(knnModel, newdata = testData$x)
modelstats <- rbind(modelstats, postResample(pred = knnPred, obs = testData$y))
postResample(pred = knnPred, obs = testData$y)
```

```
##      RMSE  Rsquared      MAE
## 3.2040595 0.6819919 2.5683461
```

MARS

```

marsGrid <- expand.grid(.degree = 1:2, .nprune = 2:38)
marsModel <- train(x = trainingData$x, y = trainingData$y, method = "earth",
                  tuneGrid = marsGrid, trControl = trainControl(method = "cv"))
marsModel

```

```


## Multivariate Adaptive Regression Spline
##
## 200 samples
## 10 predictors
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...
## Resampling results across tuning parameters:
##
## degree nprune RMSE Rsquared MAE
## 1 2 4.287738 0.2766228 3.541096
## 1 3 3.735295 0.4579537 3.037607
## 1 4 2.621465 0.7348671 2.115178
## 1 5 2.386064 0.7837287 1.928445
## 1 6 2.213116 0.8247552 1.769175
## 1 7 1.814470 0.8776482 1.436171
## 1 8 1.715089 0.8867568 1.344378
## 1 9 1.651983 0.8946079 1.296225
## 1 10 1.666730 0.8896019 1.324865
## 1 11 1.626214 0.8938909 1.285840
## 1 12 1.609580 0.8962215 1.284743
## 1 13 1.639251 0.8929914 1.297566
## 1 14 1.647570 0.8916943 1.301484
## 1 15 1.647570 0.8916943 1.301484
## 1 16 1.647570 0.8916943 1.301484
## 1 17 1.647570 0.8916943 1.301484
## 1 18 1.647570 0.8916943 1.301484
## 1 19 1.647570 0.8916943 1.301484
## 1 20 1.647570 0.8916943 1.301484
## 1 21 1.647570 0.8916943 1.301484
## 1 22 1.647570 0.8916943 1.301484
## 1 23 1.647570 0.8916943 1.301484
## 1 24 1.647570 0.8916943 1.301484
## 1 25 1.647570 0.8916943 1.301484
## 1 26 1.647570 0.8916943 1.301484
## 1 27 1.647570 0.8916943 1.301484
## 1 28 1.647570 0.8916943 1.301484
## 1 29 1.647570 0.8916943 1.301484
## 1 30 1.647570 0.8916943 1.301484
## 1 31 1.647570 0.8916943 1.301484
## 1 32 1.647570 0.8916943 1.301484
## 1 33 1.647570 0.8916943 1.301484
## 1 34 1.647570 0.8916943 1.301484
## 1 35 1.647570 0.8916943 1.301484
## 1 36 1.647570 0.8916943 1.301484
## 1 37 1.647570 0.8916943 1.301484
## 1 38 1.647570 0.8916943 1.301484
## 2 2 4.287738 0.2766228 3.541096

```

```

##      2      3      3.735295  0.4579537  3.037607
##      2      4      2.621465  0.7348671  2.115178
##      2      5      2.336068  0.7905828  1.886709
##      2      6      2.235016  0.8153197  1.740073
##      2      7      1.805188  0.8775674  1.440077
##      2      8      1.700080  0.8929903  1.314594
##      2      9      1.436055  0.9190680  1.136430
##      2     10      1.454841  0.9172529  1.143233
##      2     11      1.350271  0.9318665  1.047642
##      2     12      1.302893  0.9353308  1.032536
##      2     13      1.334764  0.9336181  1.040403
##      2     14      1.337295  0.9323946  1.049709
##      2     15      1.318129  0.9361413  1.047690
##      2     16      1.286297  0.9375534  1.016883
##      2     17      1.324754  0.9350295  1.038155
##      2     18      1.324754  0.9350295  1.038155
##      2     19      1.324754  0.9350295  1.038155
##      2     20      1.324754  0.9350295  1.038155
##      2     21      1.324754  0.9350295  1.038155
##      2     22      1.324754  0.9350295  1.038155
##      2     23      1.324754  0.9350295  1.038155
##      2     24      1.324754  0.9350295  1.038155
##      2     25      1.324754  0.9350295  1.038155
##      2     26      1.324754  0.9350295  1.038155
##      2     27      1.324754  0.9350295  1.038155
##      2     28      1.324754  0.9350295  1.038155
##      2     29      1.324754  0.9350295  1.038155
##      2     30      1.324754  0.9350295  1.038155
##      2     31      1.324754  0.9350295  1.038155
##      2     32      1.324754  0.9350295  1.038155
##      2     33      1.324754  0.9350295  1.038155
##      2     34      1.324754  0.9350295  1.038155
##      2     35      1.324754  0.9350295  1.038155
##      2     36      1.324754  0.9350295  1.038155
##      2     37      1.324754  0.9350295  1.038155
##      2     38      1.324754  0.9350295  1.038155
##
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were nprune = 16 and degree = 2.
plot(marsModel)

```



Which predictors did MARS select as the informative?

```
varImp(marsModel)
```

```
## earth variable importance
##
##      Overall
## X1    100.00
## X4     85.14
## X2     69.24
## X5     49.32
## X3     40.02
## X9       0.00
## X8       0.00
## X10      0.00
## X7       0.00
## X6       0.00
```

Only X1-X5

Performance evaluation


```
marsPred <- predict(marsModel, newdata = testData$x)
modelstats <- rbind(modelstats, postResample(pred = marsPred, obs = testData$y))
postResample(pred = marsPred, obs = testData$y)
```

```
##      RMSE Rsquared      MAE
## 1.1492504 0.9471145 0.9158382
```


SVM Radial

```
svmRModel <- train(x = trainingData$x, y = trainingData$y,
  method = "svmRadial", preProcess = c("center", "scale"),
  tuneLength = 14, trControl = trainControl(method = "cv"))
svmRModel

## Support Vector Machines with Radial Basis Function Kernel
##
## 200 samples
## 10 predictors
##
## Pre-processing: centered (10), scaled (10)
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 180, 180, 180, 180, 180, 180, ...
## Resampling results across tuning parameters:
##
##   C          RMSE      Rsquared    MAE
##   0.25  2.500756  0.7945828  2.015384
##   0.50  2.250771  0.8099288  1.806911
##   1.00  2.097422  0.8277370  1.668132
##   2.00  1.966109  0.8487748  1.547347
##   4.00  1.885043  0.8587663  1.486123
##   8.00  1.847760  0.8642743  1.452813
##  16.00  1.829788  0.8682447  1.443874
##  32.00  1.829579  0.8685594  1.444677
##  64.00  1.829579  0.8685594  1.444677
## 128.00  1.829579  0.8685594  1.444677
## 256.00  1.829579  0.8685594  1.444677
## 512.00  1.829579  0.8685594  1.444677
##1024.00  1.829579  0.8685594  1.444677
##2048.00  1.829579  0.8685594  1.444677
##
## Tuning parameter 'sigma' was held constant at a value of 0.05494384
## RMSE was used to select the optimal model using the smallest value.
## The final values used for the model were sigma = 0.05494384 and C = 32.
plot(svmRModel)
```



```
svmRModel$finalModel
```

```
## Support Vector Machine object of class "ksvm"
##
## SV type: eps-svr (regression)
## parameter : epsilon = 0.1 cost C = 32
##
## Gaussian Radial Basis kernel function.
## Hyperparameter : sigma = 0.0549438430536528
##
## Number of Support Vectors : 152
##
## Objective Function Value : -90.8443
## Training error : 0.008371
```

Performance evaluation

```
svmRPred <- predict(svmRModel, newdata = testData$x)
modelstats <- rbind(modelstats, postResample(pred = svmRPred, obs = testData$y))
postResample(pred = svmRPred, obs = testData$y)
```

```
##      RMSE Rsquared      MAE
## 2.0574214 0.8283777 1.5642193
```

Which model performs the best?

```
rownames(modelstats) <- c('NNET', 'KNN', 'MARS', 'SVM')
modelstats
```

##		RMSE	Rsquared	MAE
##	NNET	1.926316	0.8528292	1.4630365
##	KNN	3.204059	0.6819919	2.5683461
##	MARS	1.149250	0.9471145	0.9158382
##	SVM	2.057421	0.8283777	1.5642193

MARS model gives the best performance. Also it selects only informative predictors (X1-X5)